


Word Classes and Part-of-Speech (POS) Tagging

Based on slides from Julia Hirschberg -
www.cs.columbia.edu/~julia

Defining POS Tagging

- The process of assigning a part-of-speech or lexical class marker to each word in a corpus:



Applications for POS Tagging

- Speech synthesis pronunciation
 - *Lead* *Lead*
 - *INsult* *inSULT*
 - *OBject* *obJECT*
 - *OVERflow* *overFLOW*
 - *DIScount* *disCOUNT*
 - *CONtent* *conTENT*
- Parsing: e.g. *Time flies like an arrow*
 - Is *flies* an N or V?
- Word prediction in speech recognition
 - Possessive pronouns (*my, your, her*) are likely to be followed by nouns
 - Personal pronouns (*I, you, he*) are likely to be followed by verbs
- Machine Translation

Tag Ambiguity

- Words often have more than one POS: *back*
 - The *back* door = JJ
 - On my *back* = NN
 - Win the voters *back* = RB
 - Promised to *back* the bill = VB
- The POS tagging problem is *to determine the POS tag for a particular instance of a word*

Tagging Whole Sentences with POS is Hard

- Ambiguous POS contexts
 - E.g., Time flies like an arrow.
- Possible POS assignments
 - Time/[V,N] flies/[V,N] like/[V,Prep] an/Det arrow/N
 - Time/N flies/V like/Prep an/Det arrow/N
 - Time/V flies/N like/Prep an/Det arrow/N
 - Time/N flies/N like/V an/Det arrow/N
 -

Some Ways to do POS Tagging

- Rule-based tagging
 - E.g. **EnCG ENGTWOL tagger**
- Transformation-based tagging
 - Learned rules (statistical and linguistic)
 - E.g., **Brill tagger**
- Stochastic, or, Probabilistic tagging
 - **HMM (Hidden Markov Model) tagging**

POS Tagging as a Sequence ID Task

- Given a sentence (a sequence of words, or **observations**)
 - Secretariat is expected to race tomorrow
- What is the best *sequence of tags* which corresponds to this *sequence of observations*?
- Bayesian approach:
 - Consider all possible sequences of tags
 - Choose the tag sequence $t_1 \dots t_n$ which is most probable given the observation sequence of words $w_1 \dots w_n$

POS Tagging Equation

$$\hat{t}_1^n = \operatorname{argmax}_{t_1^n} P(t_1^n | w_1^n) \approx \operatorname{argmax}_{t_1^n} \prod_{i=1}^n P(w_i | t_i) P(t_i | t_{i-1})$$

Now we have two probabilities to calculate:

- Probability of a word occurring given its tag
- Probability of a tag occurring given a previous tag
- We can calculate each of these from a POS-tagged corpus

Tag Transition Probabilities $P(t_i|t_{i-1})$

- Determiners likely to precede adjs and nouns but unlikely to follow adjs
 - The/DT red/JJ hat/NN; *Red/JJ the/DT hat/NN
 - So we expect $P(NN|DT)$ and $P(JJ|DT)$ to be high but $P(DT|JJ)$ to be low
- Compute $P(NN|DT)$ by counting in a tagged corpus:

$$P(t_i|t_{i-1}) = \frac{C(t_{i-1}, t_i)}{C(t_{i-1})}$$

$$P(NN|DT) = \frac{C(DT, NN)}{C(DT)} = \frac{56,509}{116,454} = .49$$

Word Likelihood Probabilities $P(w_i|t_i)$

- VBZ (3sg pres verb) likely to be **is**
- Compute $P(\text{is}|\text{VBZ})$ by counting in a tagged corpus:

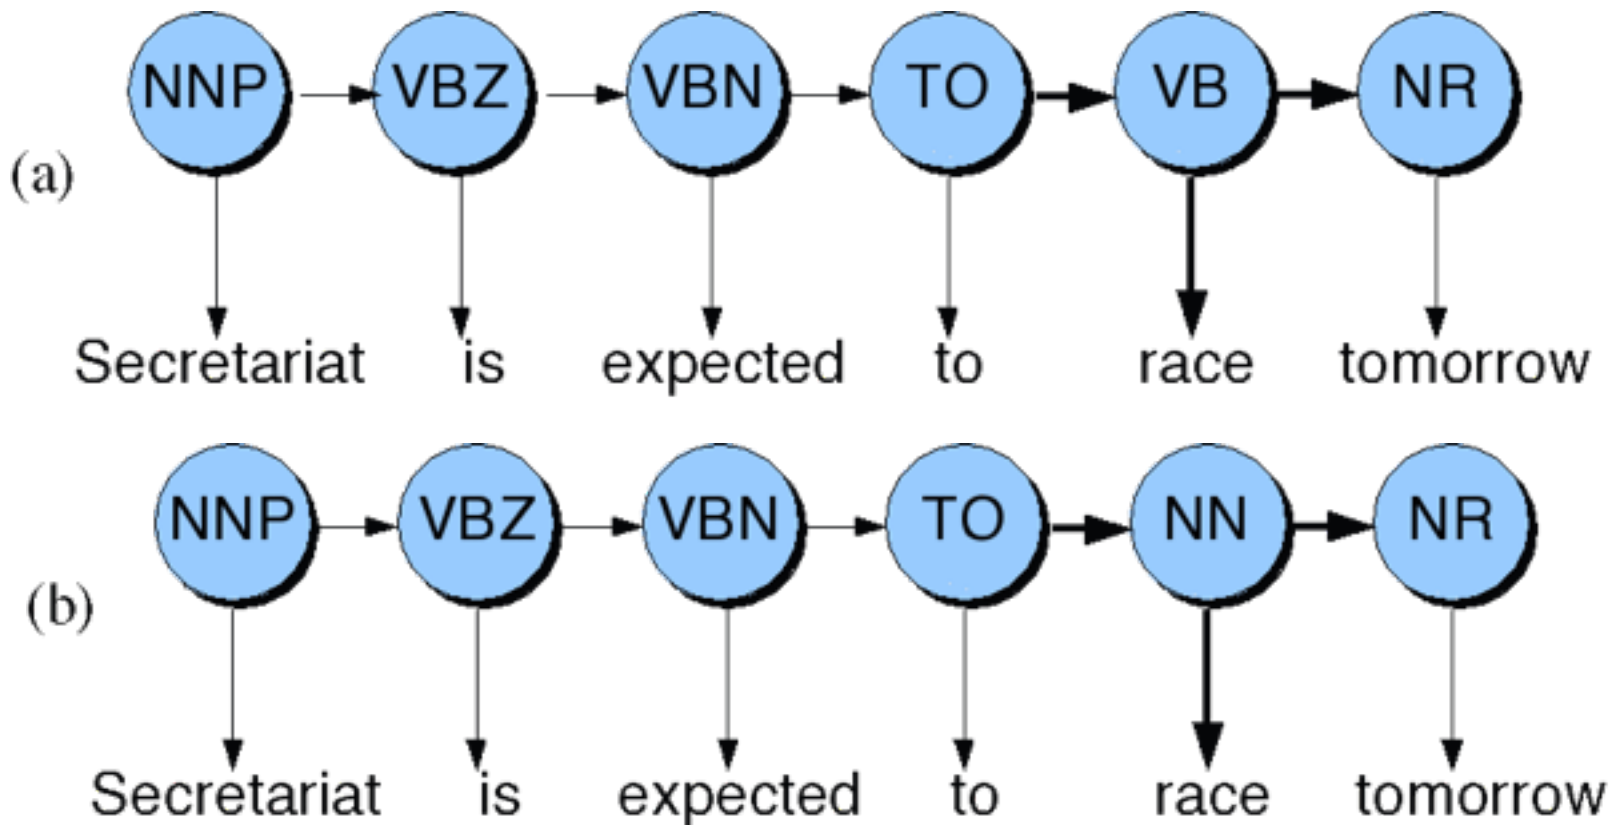
$$P(w_i|t_i) = \frac{C(t_i, w_i)}{C(t_i)}$$

$$P(\text{is}|\text{VBZ}) = \frac{C(\text{VBZ}, \text{is})}{C(\text{VBZ})} = \frac{10,073}{21,627} = .47$$

Some Data on **race**

- Secretariat/**NNP** is/**VBZ** expected/**VBN** to/**TO**
race/**VB** tomorrow/**NR**
- People/**NNS** continue/**VB** to/**TO** inquire/**VB** the/
DT reason/**NN** for/**IN** the/**DT** **race**/**NN** for/**IN**
outer/**JJ** space/**NN**
- How do we pick the right tag for **race** in new data?

Disambiguating **to race tomorrow**




$$\hat{t}_1^n = \operatorname{argmax}_{t_1^n} P(t_1^n | w_1^n) \approx \operatorname{argmax}_{t_1^n} \prod_{i=1}^n P(w_i | t_i) P(t_i | t_{i-1})$$

- $P(\text{NN}|\text{TO}) = .00047$
- $P(\text{VB}|\text{TO}) = .83$
- $P(\text{race}|\text{NN}) = .00057$
- $P(\text{race}|\text{VB}) = .00012$
- $P(\text{NR}|\text{VB}) = .0027$
- $P(\text{NR}|\text{NN}) = .0012$
- $P(\text{VB}|\text{TO})P(\text{NR}|\text{VB})P(\text{race}|\text{VB}) = .00000027$
- $P(\text{NN}|\text{TO})P(\text{NR}|\text{NN})P(\text{race}|\text{NN}) = .00000000032$
- So we (correctly) choose the verb reading

Markov Chains

- Markov Chains
 - Have **transition probabilities** like $P(t_i|t_{i-1})$
- A special case of **weighted FSTs** (FSTs which have probabilities or weights on the arcs)
- Can only represent unambiguous phenomena

A Weather Example: cold, hot, hot



Weather Markov Chain

- What is the probability of 4 consecutive rainy days?
- Sequence is rainy-rainy-rainy-rainy
- I.e., state sequence is 3-3-3-3
- $P(3,3,3,3) =$
 - $\pi_3 a_{33} a_{33} a_{33} a_{33} = 0.2 \times (0.6)^3 = 0.0432$

Weather Markov Chain

- What is the probability of 4 consecutive cold days?
- Sequence is cold-cold-cold-cold
- I.e., state sequence is 2-2-2-2
- $P(2,2,2,2) =$
 - $\pi_2 a_{22} a_{22} a_{22} a_{22} = 0.3 \times (0.5)^3 = 0.0375$

Markov Chain Defined

- A set of states $Q = q_1, q_2 \dots q_N$
- Transition probabilities
 - A set of probabilities $A = a_{01} a_{02} \dots a_{n1} \dots a_{nn}$.
 - Each a_{ij} represents the probability of transitioning from state i to state j
 - The set of these is the transition probability matrix A

$$a_{ij} = P(q_t = j \mid q_{t-1} = i) \quad 1 \leq i, j \leq N$$

$$\sum_{j=1}^N a_{ij} = 1; \quad 1 \leq i \leq N$$

- Distinguished start and end states q_0, q_F

- Can have special initial probability vector π instead of start state

$$\pi_i = P(q_1 = i) \quad 1 \leq i \leq N$$

- An initial distribution over probability of start states
- Must sum to 1

$$\sum_{j=1}^N \pi_j = 1$$

Hidden Markov Models

- Markov Chains are useful for representing problems in which
 - Observable events
 - Sequences to be labeled are unambiguous
- Problems like POS tagging are neither
- HMMs are useful for computing events we *cannot* directly observe in the world, using other events we *can* observe
 - Unobservable (Hidden): e.g., POS tags
 - Observable: e.g., words
 - We have to *learn* the relationships

Hidden Markov Models

- A set of states $Q = q_1, q_2 \dots q_N$
- Transition probabilities
 - Transition probability matrix $A = \{a_{ij}\}$

$$\sum_{j=1}^N a_{ij} = 1; \quad 1 \leq i \leq N$$

- Observations $O = o_1, o_2 \dots o_N$;
 - Each a symbol from a vocabulary $V = \{v_1, v_2, \dots, v_V\}$
- Observation likelihoods or emission probabilities
 - Output probability matrix $B = \{b_i(o_t)\}$

- Special initial probability vector π

$$\pi_i = P(q_1 = i) \quad 1 \leq i \leq N$$

- A set of legal accepting states $QA \subset Q$

First-Order HMM Assumptions

- **Markov assumption: probability of a state depends only on the state that precedes it**

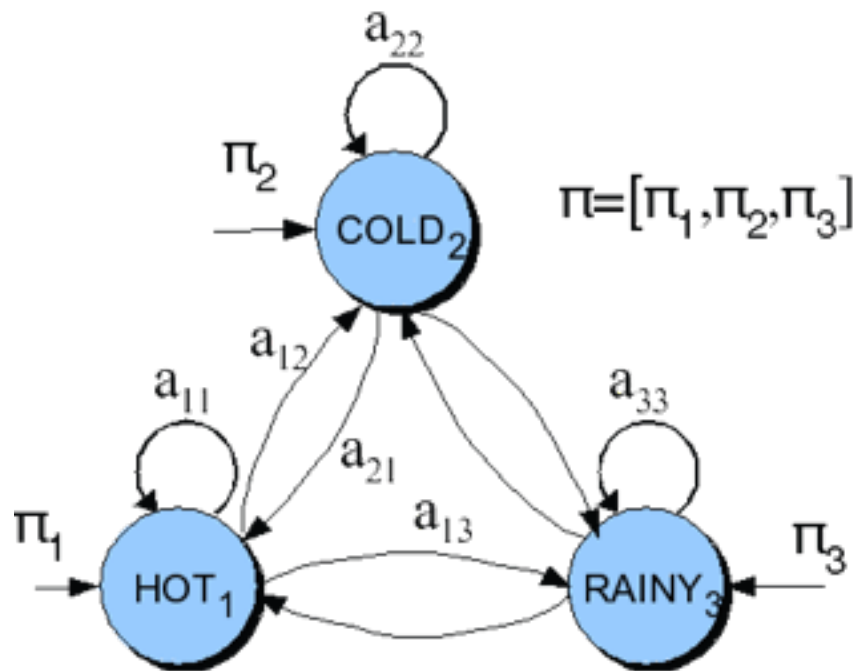
$$P(q_i | q_1 \dots q_{i-1}) = P(q_i | q_{i-1})$$

- This is the same Markov assumption we made when we decided to represent sentence probabilities as the product of bigram probabilities

- **Output-independence assumption: probability of an output observation depends only on the state that produced the observation**

$$P(o_t | O_1^{t-1}, q_1^t) = P(o_t | q_t)$$

Weather Again




Weather and Ice Cream

- You are a climatologist in the year 2799 studying global warming
- You can't find any records of the weather in Claremont for summer of 2015
- But you find Maria Klawe's diary
- Which lists how many ice-creams Maria ate every day that summer
- Your job: Determine (hidden) sequence of weather states that 'led to' Maria's (observed) ice cream behavior

Weather/Ice Cream HMM

- Hidden States: {Hot,Cold}
- Transition probabilities (A Matrix) between H and C
- Observations: {1,2,3} # of ice creams eaten per day
- Goal: Learn observation likelihoods between observations and weather states (Output Matrix B) by training HMM on aligned input streams from a training corpus
- Result: trained HMM for weather prediction given ice cream information alone

Ice Cream/Weather HMM



What can HMMs Do?

- ***Likelihood***: Given an HMM $\lambda = (A, B)$ and an observation sequence O , determine the likelihood $P(O, \lambda)$: Given # ice creams, what is the weather?
- ***Decoding***: Given an observation sequence O and an HMM $\lambda = (A, B)$, discover the best hidden state sequence Q : Given seq of ice creams, what was the most likely weather on those days?
- ***Learning***: Given an observation sequence O and the set of states in the HMM, learn the HMM parameters A and B

Decoding: The Viterbi Algorithm

- Decoding: Given an observation sequence O and an HMM $\lambda = (A, B)$, discover the *best* hidden state sequence of weather states in Q
 - E.g., Given the observations 3 – 1 – 1 and an HMM, what is the *best* (most probable) hidden weather sequence of $\{H, C\}$
- Viterbi algorithm
 - Dynamic programming algorithm
 - Uses a dynamic programming trellis to store probabilities that the HMM is in state j after seeing the first t observations, for all states j

- Value in each cell computed by taking MAX over all paths leading to this cell – i.e. best path
- Extension of a path from state i at time $t-1$ is computed by multiplying:

$$v_t(j) = \max_{1 \leq i \leq N-1} v_{t-1}(i) a_{ij} b_j(o_t)$$

$v_{t-1}(i)$ the **previous Viterbi path probability** from the previous time step
 a_{ij} the **transition probability** from previous state q_i to current state q_j
 $b_j(o_t)$ the **state observation likelihood** of the observation symbol o_t given the current state j

- Most probable path is the max over all possible previous state sequences

HMM Training: The Forward-Backward (Baum-Welch) Algorithm

- ***Learning***: Given an observation sequence O and the set of states in the HMM, learn the HMM parameters A (transition) and B (emission)
- Input: unlabeled seq of observations O and vocabulary of possible hidden states Q
 - E.g. for ice-cream weather:
 - Observations = $\{1,3,2,1,3,3,\dots\}$
 - Hidden states = $\{H,C,C,C,H,C,\dots\}$

- Intuitions
 - Iteratively re-estimate counts, starting from an initialization for A and B probabilities, e.g. all equiprobable
 - Estimate new probabilities by computing **forward probability** for an observation, dividing prob. mass among all paths contributing to it, and computing the **backward probability** from the same state
- Details: see e.g. Manning & Schütze

What is Syntax?


- ✿ The study of how the parts of an utterance are arranged in relation to one another
- ✿ Questions in syntax:
 - ✿ Do all languages behave the same way?
 - ✿ Can the structure of yet un-analyzed languages be predicted?
 - ✿ How is syntax learned by children, with little negative evidence?

Structural Descriptions

- ✱ A structure that shows word order, syntactic constituency, and labels for the constituents.
- ✱ Includes trees, bracketed structures

Structural Descriptions

Tree



Bracketed Structure

[[The [dog]] [bit [the [man]]]]

Labeled bracketed structure

[S [NP [det The] [Nom [N dog]]] [VP [V bit] [NP [Det the] [Nom [N man]]]]]

Tree Structures

- ☼ Ordered directed trees with nodes, labels, arcs
 - ☼ Preterminal Node: Node with a single leaf as its descendant
 - ☼ What are preterminal nodes in NL grammar?
 - ☼ Part of Speech Tags
 - ☼ Arc: Shows constituency relation, but untyped
 - ☼ Label: Symbol giving the category of a node

What is Syntax? (part 2)

- ☼ Set of rules by which well formed utterances are formed.
- ☼ Formalize the notion of syntax:
formal language theory

Formal Language Theory

- ✱ Natural language is rule-governed, not random
- ✱ Like compilers, machine languages
- ✱ Can construct a grammar to parse it
- ✱ Formal language theory
 - ✱ Conceptual framework for studying natural language.

Generalized Formal Grammar

- ✻ $G = \langle N, \Sigma, P, S \rangle$ where:
 - ✻ N is a set of non-terminal symbols, typically S, A, B, \dots
 - ✻ Σ is a set of terminals, typically x, y, z, \dots
 - ✻ P is a set of production rules
 - ✻ S is the starting or goal variable from N , i.e., $S \in N$

Sample Grammar

$S \rightarrow NP VP$

$NP \rightarrow Det Noun$

$NP \rightarrow ProperNoun$

$VP \rightarrow Verb$

$VP \rightarrow Verb NP$

$Det \rightarrow the \mid a \mid that$

$Noun \rightarrow lamp \mid pig \mid dirt$

$ProperNoun \rightarrow Washington \mid Sam$

$Verb \rightarrow understands \mid chases$

Washington understands Sam

Sam chases that pig

*understands Sam

Context Free Grammar

$G = \langle N, \Sigma, P, S \rangle$ where:

- ✿ N is a set of non-terminal symbols, typically S, A, B, \dots
- ✿ S is the starting or goal symbol from N , i.e., $S \in N$
- ✿ Σ is a set of terminal symbols, typically x, y, z, \dots disjoint from N
- ✿ P is a set of production rules of the form $A \rightarrow \beta$, where:
 - ✿ A is a non-terminal $A \in N$
 - ✿ β is a string of symbols from $(\Sigma \cup N)$

CFGs for Natural Language

- ✻ A nonterminal symbol labels a syntactic part (constituent):
NP, VP, PP, (Noun, Verb, Det)
- ✻ A starting symbol indicates which symbol has to come first; it labels the largest constituent or biggest part:
S, ROOT, or TOP
- ✻ A terminal symbol labels the smallest part, the actual strings of the language:
man, they, swim

CFGs for Natural Language

- ☼ Production rule (re-write rule): one symbol is rewritten (\rightarrow) as one or more others:
 $NP \rightarrow Det\ Noun$
- ☼ A production rule captures the notion of syntactic constituency.
- ☼ ‘LHS’ is used to indicate the left-hand side of the \rightarrow , and likewise for ‘RHS’.

Is this a valid CFG grammar?

$NP \rightarrow \text{Det Noun}$

$\text{Nom} \rightarrow (\text{Adj}) \text{Noun}$

$VP \rightarrow \text{VB NP}$

$\text{Det} \rightarrow \text{the} \mid \text{a}$

$\text{Noun} \rightarrow \text{rabbit} \mid \text{carrots}$

$\text{Adj} \rightarrow \text{fresh} \mid \text{crispy}$

$\text{VB} \rightarrow \text{ate} \mid \text{likes}$

the rabbit likes crispy carrots.

Treebanks

- ✿ Linguistic corpora annotated for syntactic structure.
- ✿ Imply grammars of the languages they contain
- ✿ Examples:
 - ✿ The Penn Treebank (English)
 - ✿ Penn Chinese Treebank Project
 - ✿ The Tübingen Treebank of Written German
 - ✿ Arabic Treebank
 - ✿ Korean Treebank

Some Phrases in PTB

- ☼ NP: Noun phrase
- ☼ VP: Verb phrase
- ☼ PP: Prepositional phrase
- ☼ ADJP: Adjective phrase
- ☼ ADVP: Adverb phrase
- ☼ CONJP: Multi-word conjunctions (“not only”)
- ☼ QP: quantifier phrase (inside NPs)
- ☼ 21 Total

Clauses in PTB

- ✿ S: declaratives, passives, imperatives, questions with declarative order, (embedded) infinitive clauses, gerund classes
- ✿ SINV: Inverted clauses
- ✿ SBAR: Relative and subordinate clauses
- ✿ SBARQ :Wh-questions
- ✿ SQ: Y/N-questions, inside SBARQ
- ✿ S-CLF : It-cleft clauses
- ✿ FRAG: Stand-alone clauses, phrases without a predicate argument structure

Rules in Treebanks

- ✱ Lots of them! 17,000 in PTB
 - ✱ Most very flat
 - ✱ Many tailored to single sentences
 - ✱ Number grows linearly with corpus
- ✱ Largest number: S, NP, VP

Two Goal of Parsing

Analyze input strings to assign proper structures

- ✱ For input A , grammar G :
 - ✱ Assign zero or more parse tree(s) T :
 - ✱ Cover all and only the elements of A
 - ✱ Root of T is S (the start symbol of G)
 - ✱ Do not necessarily pick one (or correct) analysis

Two Goal of Parsing

Recognition

- ✱ Subtask of parsing
- ✱ For input A , grammar G
 - ✱ Is A in the language defined by G ?

Questions for Parsing

- ✿ Is this sentence in the language?
 - ✿ FSAs accept the regular languages defined by automaton
 - ✿ Parsers accept language defined by CFG
- ✿ What is the syntactic structure of this sentence?
 - ✿ Syntactic parse provides framework for semantic analysis
 - ✿ What is the subject?
 - ✿ Useful for e.g. question answering

Parsing as Search

- ✱ Search through possible parse trees
- ✱ Want one (or more) that derive input
- ✱ Formally, search problems are defined by:
 - ✱ Start state S ,
 - ✱ Goal state G ,
 - ✱ Successor Function:
Transitions between states,
 - ✱ Path cost function

One Model of Parsing as Search

- ☼ Start State:
 - ☼ Start Symbol from grammar
- ☼ Goal test:
 - ☼ Does parse tree cover all and only input?
- ☼ Successor function:
 - ☼ Expand a non-terminal using production in grammar where non-terminal is LHS of grammar
- ☼ Path cost:
 - ☼ We'll ignore here

One Model of Parsing as Search

- ✱ Node: Partial solution to search problem:
 - ✱ Partial parse
- ✱ Search start node: Initial State
 - ✱ Input string
 - ✱ Start symbol of CFG
- ✱ Goal node:
 - ✱ Full parse tree: covering all and only input, rooted at S

Parse Search Strategies

- ✱ Two constraints:
 - ✱ Must start with the start symbol
 - ✱ Must cover exactly the input string
- ✱ Correspond to main parsing search strategies
 - ✱ Top-down search (Goal-directed search)
 - ✱ Bottom-up search (Data-driven search)

Parse Search Strategies

	Breadth-First	Depth-First
Top-Down		
Bottom-Up		